

# Linear Regression and Regularization

by Dane Brown

**“Life is ten percent what you experience and ninety percent how you respond to it.”**

# Scoring a Regression Model

- *mean\_squared\_error*: squared error between **predicted** and **true value** for every data point in the training set, averaged across all data points.
- *explained\_variance\_score*: the degree a model can explain the variation or dispersion of the test data. Measured using the correlation coefficient.
- *r2\_score*: The  $R^2$  score is closely related to the explained variance score, but uses an unbiased variance estimation. It is also known as the coefficient of determination.

# cvML Methodology

- **Initialization:** Call the *cv* or *scikit* model by name to create an empty instance of the model
- **Set parameters:** default
- **Train the model:** *train* or *fit* is used to fit the model to some data
- **Predict new labels:** use *predict*, to guess the labels of new (**unseen**) data
- **Score the model:** refer to slide 10 & 13: works for both *cv* or *scikit*

# Linear Regression

- Continuously predict new outcomes
- Describe a target variable with a linear combination of features
- Dataset: Boston house pricing
  - simply predict housing prices
  - no classifying of labels (into classes)
  - $\hat{y} = w_1 f_1 + w_2 f_2$
  - if  $f_1$  and  $f_2$  are today's price for two houses
  - train  $w_1$  and  $w_2$  to learn tomorrow's price
- OpenCV does not offer any good implementation of linear regression...

# Linear Regression

Check it out -> [CV\\_ML](#)

- The goal is to predict the value of homes in several Boston neighborhoods in the 1970s, using information such as
  - crime rate
  - property tax rate
  - distance to employment centers
  - highway accessibility
  - etc.

# Linear Regression Results


The model is noticeably off at certain points...



# Linear Regression Results

- The model tends to be off the most at peaks
  - i.e. really high or really low housing prices
  - peak values of data points 12, 18, and 42.

# Linear Regression Results





# Linear Regression Results

- A perfect model has all data points on the dashed diagonal, since **y\_pred** would equal **y\_true**
- Deviations from the diagonal indicate model errors: the model was not able to explain some variance in the data .
- R-squared shows that the model explained 76% of the scatter in the data, with an MSE of 15.011.
- ML project: These are the kind of numbers to use in comparisons and results chapter of your thesis!

# Overfitting

- An algorithm might work really well on the training set, but poorly on unseen data
- This means the model does not generalize well
- This is especially common in tree-based classifiers
- This can also happen for other classifiers, including **regression** problems

# Overfitting: How to Reduce it

- Use Regularization
- L1 norm of the Manhattan distance:
  - This adds a term to the scoring function that is proportional to the sum of all absolute weight values
  - Also known as **Lasso** regression
- L2 norm of the Euclidean distance:
  - This adds a term to the scoring function that is proportional to the sum of all squared weight values.
  - Removes strong outliers in the weight vector much more than the L1 norm
  - Also known as **Ridge** regression

# Overfitting: How to Reduce it

- Original Linear regression:
  - `linreg = linear_model.LinearRegression()`
- Lasso regression
  - `linreg = linear_model.Lasso()`
- Ridge regression
  - `linreg = linear_model.RidgeRegression()`

# Lasso vs. Ridge

Neither are better in general

Each tend to do well under conditions:

- **Lasso**
  - a small number of significant parameters and the others are close to zero significance
  - i.e. when few predictors influence the response
- **Ridge**
  - many significant parameters of similar value
  - i.e. when most predictors impact the response
- The boston dataset has many predictor variables